

# Algorithm 941: htucker—A MATLAB Toolbox for Tensors in Hierarchical Tucker Format

DANIEL KRESSNER and CHRISTINE TOBLER, EPF Lausanne

The hierarchical Tucker format is a storage-efficient scheme to approximate and represent tensors of possibly high order. This article presents a MATLAB toolbox, along with the underlying methodology and algorithms, which provides a convenient way to work with this format. The toolbox not only allows for the efficient storage and manipulation of tensors in hierarchical Tucker format but also offers a set of tools for the development of higher-level algorithms. Several examples for the use of the toolbox are given.

Categories and Subject Descriptors: G.1.3 [Numerical Analysis]: Numerical Linear Algebra—*Sparse, structured, and very large systems (direct and iterative methods)*; *Singular value decomposition*; G.4 [Mathematical Software]: *Algorithm design and analysis*

General Terms: Algorithms

Additional Key Words and Phrases: Higher-order tensors, multilinear algebra, tensor networks, MATLAB

## ACM Reference Format:

Kressner, D. and Tobler, C. 2014. Algorithm 941: htucker—A MATLAB toolbox for tensors in hierarchical Tucker format. *ACM Trans. Math. Softw.* 40, 3, Article 22 (April 2014), 22 pages.

DOI: <http://dx.doi.org/10.1145/2538688>

## 1. INTRODUCTION

A tensor  $\mathcal{X} \in \mathbb{C}^{n_1 \times n_2 \times \dots \times n_d}$  with  $n_1, \dots, n_d \in \mathbb{N}$  is a  $d$ -dimensional array with entries  $\mathcal{X}_{i_1 i_2 \dots i_d} \in \mathbb{C}$ . Usually,  $d$  is called the *order* of the tensor and the focus of this article is on tensors of higher order, say,  $d = 5$  or  $d = 10$  or even  $d = 100$ . A typical scenario is that  $\mathcal{X}$  represents a  $d$ -variate function  $f : [0, 1]^d \rightarrow \mathbb{C}$  sampled on a tensor grid or approximated in a tensorized basis.

It is in general impossible to store a higher-order tensor explicitly, simply because the number of entries grows exponentially with  $d$ . Various data-sparse formats have been developed to address this issue. Depending on the application, these formats may allow for the approximate representation and manipulation of a tensor under dramatically reduced storage and computing requirements. For example, consider the approximation of  $\mathcal{X}$  by a rank-1 tensor

$$\text{vec}(\mathcal{X}) \approx u_d \otimes u_{d-1} \otimes \dots \otimes u_1, \quad u_1 \in \mathbb{C}^{n_1}, \dots, u_d \in \mathbb{C}^{n_d}, \quad (1)$$

where  $\text{vec}$  stacks the entries of a tensor in reverse lexicographical order into a long column vector and  $\otimes$  denotes the standard Kronecker product. Then, instead of the  $n_1 \cdot n_2 \cdot \dots \cdot n_d$  entries of  $\mathcal{X}$ , only the  $n_1 + n_2 + \dots + n_d$  entries of  $u_1, \dots, u_d$  need to be stored. On function level, this corresponds to an approximation of  $f$  by a separable function.

---

This work was supported by the FNS research module *Preconditioned methods for large-scale model reduction* within the FNS ProDoc *Efficient Numerical Methods for Partial Differential Equations*.

Authors' address: D. Kressner and C. Tobler (corresponding author), Chair of Numerical Algorithms and HPC, MATHICSE, EPF Lausanne, CH-1915 Lausanne, Switzerland; email: [christine.tobler@gmx.ch](mailto:christine.tobler@gmx.ch).

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

© 2014 ACM 0098-3500/2014/04-ART22 \$15.00

DOI: <http://dx.doi.org/10.1145/2538688>

A typical application we have in mind is when  $\mathcal{X}$  arises from the discretization of a high-dimensional or parameter-dependent partial differential equation and is only given implicitly as the solution to a typically huge (non)linear system or eigenvalue problem. There are two quite different strategies to employ a data-sparse format for the solution of such problems. The more straightforward one is to apply a standard iterative method, such as a conjugate gradient method, and approximate each iterate in the data-sparse format. For this purpose, it is desirable to keep the approximation error negligible; otherwise the accuracy and convergence of the method may be compromised. Examples for this strategy can be found in Ballani and Grasedyck [2013], Hackbusch et al. [2012], Khoromskij and Oseledets [2010], Khoromskij and Schwab [2011], and Kressner and Tobler [2011]. The second strategy is to reformulate the problem at hand as an optimization problem with the admissible set of solutions restricted to data-sparse tensors; see Espig et al. [2012a], Holtz et al. [2012], Huckle et al. [2013], Koch and Lubich [2010], Oseledets and Khoromskij [2010], and Schollwöck [2011] and references therein. Beyond these two main strategies, there exist further approaches tailored to particularly structured problems, for example, Grasedyck [2004] and Kressner and Tobler [2010]. While the mathematical understanding is still somewhat limited, there is strong numerical evidence that such data-sparse algorithms can handle a wide variety of problems that are far from tractable by classical numerical methods.

In most applications, it is unlikely that a rank-1 representation (1) yields a satisfactory approximation error. This can be improved by considering the more general CP (Canonical Polyadic) decomposition [Carroll and Chang 1970; Harshman 1970]

$$\text{vec}(\mathcal{X}) \approx \sum_{j=1}^R u_d^{(j)} \otimes u_{d-1}^{(j)} \otimes \cdots \otimes u_1^{(j)}, \quad u_1^{(j)} \in \mathbb{C}^{n_1}, \dots, u_d^{(j)} \in \mathbb{C}^{n_d}, \quad (2)$$

which still requires little memory, provided that  $R$  does not become too large. Unfortunately, developing a robust and efficient algorithm for this format, which yields an approximation to any desirable accuracy, remains a subtle problem; see Acar et al. [2011] and Espig and Hackbusch [2012] for recent progress. This problem is much less subtle for the Tucker decomposition [Tucker 1966],

$$\text{vec}(\mathcal{X}) \approx (U_d \otimes U_{d-1} \otimes \cdots \otimes U_1) \text{vec}(\mathcal{C}), \quad U_1 \in \mathbb{C}^{n_1 \times r_1}, \dots, U_d \in \mathbb{C}^{n_d \times r_d}, \quad (3)$$

with the so-called core tensor  $\mathcal{C} \in \mathbb{C}^{r_1 \times r_2 \times \cdots \times r_d}$ . The HOSVD (Higher-Order SVD) [De Lathauwer et al. 2000] provides a simple, nearly optimal solution to the approximation problem (3). However, the need for storing  $\mathcal{C}$  still results in memory requirements that grow exponentially with  $d$ .

Motivated by the limitations of the two classical decompositions (2) and (3), various other decompositions have been developed in the numerical analysis community with the aim of combining the advantages of both. This includes the tensor train decomposition [Khoromskij and Oseledets 2011], the closely related but somewhat more general HTD (hierarchical Tucker decomposition) [Grasedyck 2010b; Hackbusch and Kühn 2009], and even more general tensor networks [Espig et al. 2011]. In the computational physics community, matrix product states and tensor networks play a central role in DMRG (density matrix renormalization group) techniques for computing ground states of quantum many-body systems; see Schollwöck [2011] for an introduction. In the case of tensors that contain values of a multivariate function, additional properties inherited from the regularity of the function can in certain cases lead to a significantly higher compression in HTD (and related decompositions) compared to Tucker; see, for example, Hackbusch [2012] for such a discussion.

Existing MATLAB toolboxes for working with low-rank tensor formats are the N-way toolbox by Andersson and Bro [2000], the Tensor Toolbox by Bader and Kolda [2006, 2007], as well as the TT-Toolbox by Oseledets [2011a, 2011b]. TensorCalculus [Espig et al. 2012b] is a C++ library for more general tensor networks. In computational physics, a number of related software packages have been developed in the context of DMRG techniques for simulating quantum networks; see, for example, Bauer et al. [2011].

The main goal of this article is to provide a convenient framework for the development and implementation of algorithms based on the hierarchical Tucker decomposition. In particular, methods for the truncation of a given tensor to HTD, namely, the approximation by a low-rank tensor in HTD, are implemented and discussed in Section 6. A set of advanced tools allows for the development of higher-level algorithms. Moreover, this work contains a number of new contributions.

- Efficient algorithms for HTD tensor-tensor contraction, including the inner product between two tensors, are given in Sections 5.3 and 5.4.
- A new variant of truncation to HTD without initial orthogonalization is presented in Section 6.2.2 and demonstrated to result in an efficient and numerically robust way for adding tensors in Section 6.3.
- New algorithms for exact and approximate elementwise multiplication of tensors in HTD are presented in Section 7.
- A framework for representing linear operators in HTD is presented in Section 8 and an explicit representation for a discretized Laplace operator has been given.

All our algorithms are mainly based on calls to level-3 BLAS and LAPACK functionality.

The rest of this article is organized as follows. Section 2 introduces basic tools for working with tensors as well as the HTD. In Section 3, we describe the basic functionality and data structures of our MATLAB toolbox `htucker`. Section 4 is concerned with basic operations on tensors in HTD, such as  $\mu$ -mode matrix products and orthogonalization, and their implementation in `htucker`. Tensor-tensor contractions are discussed in Section 5, including, for example, inner products. In Section 6, we present several methods for approximating a tensor either given explicitly or given in HTD by a tensor in HTD (of lower rank). Section 7 is concerned with elementwise multiplication, and Section 8 with the representation of linear operators on tensors in HTD.

Examples for working with `htucker` and a list of its complete functionality can be found in the preprint [Kressner and Tobler 2012], an extended version of this manuscript.

## 2. PRELIMINARIES

This section summarizes the mathematical foundation of the hierarchical Tucker decomposition (HTD) and our MATLAB toolbox. Necessary tensor concepts will be briefly introduced. We refer the reader to the survey paper [Kolda and Bader 2009] and to the book [Hackbusch 2012] for a more comprehensive introduction.

### 2.1. Matricization and HOSVD

To understand the principles behind HTD, it is helpful to recall the matricization and HOSVD of tensors. A tensor  $\mathcal{X} \in \mathbb{C}^{n_1 \times n_2 \times \dots \times n_d}$  has  $d$  different *modes*  $1, \dots, d$ . Consider a splitting of these modes into two disjoint sets:  $\{1, \dots, d\} = t \cup s$  with  $t = \{t_1, \dots, t_k\}$  and  $s = \{s_1, \dots, s_{d-k}\}$ . Then the corresponding *matricization* of these modes is

obtained by merging the first group into row indices and the second group into column indices

$$X^{(t)} \in \mathbb{C}^{(n_{t_1} \cdots n_{t_k}) \times (n_{s_1} \cdots n_{s_{d-k}})} \quad \text{with} \quad \left( X^{(t)} \right)_{(i_{t_1}, \dots, i_{t_k}), (i_{s_1}, \dots, i_{s_{d-k}})} := \mathcal{X}_{i_1, \dots, i_d}$$

for any indices  $i_1, \dots, i_d$  in the multi-index set  $\{1, \dots, n_1\} \times \cdots \times \{1, \dots, n_d\}$ . Of course, the order in which the indices are merged is important. In the following, we assume reverse lexicographical order of the multi-indices but any other consistently employed order would be suitable. Note that the matricization is *not* independent of the ordering of the modes  $t_1, \dots, t_k$  and  $s_1, \dots, s_{d-k}$  in the sets  $t$  and  $s$ , respectively. If not noted otherwise, we assume the modes to be in increasing order.

As a special case, consider the so-called  $\mu$ -mode matricization

$$X^{(\mu)} \in \mathbb{C}^{n_\mu \times (n_1 \cdots n_{\mu-1} n_{\mu+1} \cdots n_d)}, \quad \mu = 1, \dots, d.$$

Then the tuple  $(r_1, \dots, r_d)$  with  $r_\mu = \text{rank}(X^{(\mu)})$  is called the *multilinear rank* of  $\mathcal{X}$ . To obtain an approximation of lower multilinear rank  $(\tilde{r}_1, \dots, \tilde{r}_d)$ , with  $\tilde{r}_\mu \leq r_\mu \leq n_\mu$ , we let  $U_\mu \in \mathbb{C}^{n_\mu \times \tilde{r}_\mu}$  contain the  $\tilde{r}_\mu$  dominant left singular vectors of  $X^{(\mu)}$ , which can be obtained, for example, from a truncated SVD  $X^{(\mu)} \approx U_\mu \Sigma_\mu V_\mu^H$ . Then the HOSVD takes the form of a Tucker decomposition

$$\text{vec}(\mathcal{X}) \approx \text{vec}(\tilde{\mathcal{X}}) := (U_d \otimes \cdots \otimes U_1) \text{vec}(\mathcal{C}), \quad (4)$$

with the core tensor

$$\text{vec}(\mathcal{C}) := (U_d^H \otimes \cdots \otimes U_1^H) \text{vec}(\mathcal{X}) \in \mathbb{C}^{\tilde{r}_1 \times \cdots \times \tilde{r}_d}.$$

This choice of  $\mathcal{C}$  minimizes  $\|\mathcal{X} - \tilde{\mathcal{X}}\|_2$  for given  $U_1, \dots, U_d$  with orthonormal columns. Here and in the following,  $\|\mathcal{Y}\|_2$  denotes the Euclidean norm of the vectorization:  $\|\mathcal{Y}\|_2 = \|\text{vec}(\mathcal{Y})\|_2$ . An important feature of the HOSVD, it can be shown [De Lathauwer et al. 2000] that the obtained approximation is nearly optimal among *all* tensors of multilinear rank  $(\tilde{r}_1, \dots, \tilde{r}_d)$  or lower:

$$\|\mathcal{X} - \tilde{\mathcal{X}}\|_2 \leq \sqrt{d} \cdot \inf \{ \|\mathcal{X} - \mathcal{Y}\|_2 : \text{rank}(\mathcal{Y}^{(\mu)}) \leq \tilde{r}_\mu, \mu = 1, \dots, d \}. \quad (5)$$

## 2.2. The Hierarchical Tucker Decomposition (HTD)

In contrast to the Tucker decomposition, HTD employs a hierarchy of matricizations, motivated by the following nestedness property. Note that our use of the symbol  $\subset$  for set inclusion allows for equal sets.

**LEMMA 2.1** [GRASEDYCK 2010A, LEMMA 17]. *Let  $\mathcal{X} \in \mathbb{C}^{n_1 \times \cdots \times n_d}$  and  $t = t_l \cup t_r$  for  $t_l = \{i_l, i_l + 1, \dots, i_m\}$  and  $t_r = \{i_m + 1, \dots, i_r\}$ . Then  $\text{span}(X^{(t)}) \subset \text{span}(X^{(t_r)} \otimes X^{(t_l)})$ .*

**PROOF.** Any column of  $X^{(t)} = X^{(i_l, \dots, i_r)}$  can be considered as the vectorization of a tensor  $\mathcal{C} \in \mathbb{C}^{n_{i_l} \times \cdots \times n_{i_r}}$ . The columns of the matricization  $C^{(t_l)}$  are clearly contained in  $\text{span}(X^{(t_l)})$  and hence


$$C^{(t_l)} = X^{(t_l)} (X^{(t_l)})^+ C^{(t_l)},$$

where  $M^+$  denotes the Moore-Penrose pseudoinverse of a matrix  $M$ . Analogously,

$$(C^{(t_l)})^T = C^{(t_r)} = X^{(t_r)} (X^{(t_r)})^+ C^{(t_r)}.$$

These two relations imply

$$C^{(t_l)} = X^{(t_l)} \underbrace{\left( (X^{(t_l)})^+ C^{(t_l)} (X^{(t_r)})^{+T} \right)}_{=:V} (X^{(t_r)})^T \Rightarrow \text{vec}(\mathcal{C}) = (X^{(t_r)} \otimes X^{(t_l)}) \text{vec}(V). \quad \square$$

Fig. 1. Illustration of the HTD (7) for  $d = 4$ .

Given any bases  $U_t, U_{t_l}, U_{t_r}$  for the column spaces of  $X^{(t)}, X^{(t_l)}, X^{(t_r)}$ , the result of Lemma 2.1 implies the existence of a so-called transfer matrix  $B_t$  such that

$$U_t = (U_{t_r} \otimes U_{t_l})B_t, \quad B_t \in \mathbb{C}^{r_{t_l} r_{t_r} \times r_t}, \quad (6)$$

where  $r_t, r_{t_l}, r_{t_r}$  denote the ranks of the corresponding matricizations. Applying this relation recursively, until  $t_l$  and  $t_r$  become singletons, leads to the HTD.

*Example 2.2.* Repeated application of (6) for  $d = 4$

$$\begin{aligned} \text{vec}(\mathcal{X}) &= X^{(\{1,2,3,4\})} = (U_{34} \otimes U_{12})B_{1234} \\ U_{12} &= (U_2 \otimes U_1)B_{12} \\ U_{34} &= (U_4 \otimes U_3)B_{34} \\ \Rightarrow \text{vec}(\mathcal{X}) &= (U_4 \otimes U_3 \otimes U_2 \otimes U_1)(B_{34} \otimes B_{12})B_{1234}. \end{aligned} \quad (7)$$

It is often advantageous to reshape the transfer matrices:

$$\begin{aligned} B_{1234} \in \mathbb{C}^{r_{12}r_{34} \times 1} &\Rightarrow \mathcal{B}_{1234} \in \mathbb{C}^{r_{12} \times r_{34} \times 1}, \\ B_{12} \in \mathbb{C}^{r_1 r_2 \times r_{12}} &\Rightarrow \mathcal{B}_{12} \in \mathbb{C}^{r_1 \times r_2 \times r_{12}}, \\ B_{34} \in \mathbb{C}^{r_3 r_4 \times r_{34}} &\Rightarrow \mathcal{B}_{34} \in \mathbb{C}^{r_3 \times r_4 \times r_{34}}. \end{aligned}$$

To avoid cluttering the notation we write  $B_{12}$  instead of  $B_{\{1,2\}}$ ,  $U_{12}$  instead of  $U_{\{1,2\}}$ ,  $r_{12}$  instead of  $r_{\{1,2\}}$ , and so on.


An illustration of the hierarchical structure and the data to be stored for the HTD (7) is given in Figure 1.

The general construction of an HTD requires a hierarchical splitting of the modes  $1, \dots, d$ .

*Definition 2.3.* A binary tree  $\mathcal{T}$  with each node represented by a subset of  $\{1, \dots, d\}$  is called a *dimension tree* if the root node is  $\{1, \dots, d\}$ , each leaf node is a singleton, and each parent node is the disjoint union of its two children. In the following, we denote

$$\begin{aligned} \mathcal{L}(\mathcal{T}) &\text{ set of all leaf nodes;} \\ \mathcal{N}(\mathcal{T}) &\text{ set of all non-leaf nodes, } \mathcal{N}(\mathcal{T}) = \mathcal{T} \setminus \mathcal{L}(\mathcal{T}). \end{aligned}$$

*Remark 2.4.* For convenience of notation, we impose the following additional assumption on the left and right children  $t_l, t_r$  of a node  $t$  in the dimension tree: Each element of  $t_l$  is smaller than any element of  $t_r$ . Note that this assumption can always be satisfied by an appropriate reordering of the modes.

Fig. 2. A dimension tree for  $d = 7$ .

It is not hard to see that the number of non-leaf nodes is always  $d - 1$ . An example of a dimension tree is given in Figure 2.

Having prescribed a maximal rank  $k_t$  for each node  $t \in \mathcal{T}$ , the set of *hierarchical Tucker tensors of hierarchical rank at most  $(k_t)_{t \in \mathcal{T}}$*  is defined as

$$\mathcal{H}\text{-Tucker}((k_t)_{t \in \mathcal{T}}) = \left\{ \mathcal{X} \in \mathbb{C}^{n_1 \times \dots \times n_d} : \text{rank}(X^{(t)}) \leq k_t \text{ for all } t \in \mathcal{T} \right\}.$$

Such a hierarchical Tucker tensor  $\mathcal{X}$  is stored in the *hierarchical Tucker format* as follows. At each leaf node  $\{\mu\}$  a basis  $U_\mu \in \mathbb{C}^{n_\mu \times r_\mu}$ , where  $r_\mu := \text{rank}(X^{(\mu)}) \leq k_\mu$ , is stored. At each parent node  $t$  with children  $t_l$  and  $t_r$ , the third-order transfer tensor  $B_t \in \mathbb{C}^{r_{t_l} \times r_{t_r} \times r_t}$  satisfying (6) with  $B_t \equiv B_t^{((1,2))}$  is stored. Equivalently, (6) can be written as

$$(U_t)_{:,q} = \sum_{i=1}^{r_{t_l}} \sum_{j=1}^{r_{t_r}} \left( (U_{t_r})_{:,j} \otimes (U_{t_l})_{:,i} \right) (B_t)_{i,j,q}, \quad q = 1, 2, \dots, r_t. \quad (8)$$

The HTD for  $\mathcal{X}$  is obtained by recursively inserting (6) as illustrated in Example 2.2.

In summary, the hierarchical Tucker format is represented by  $d$  matrices  $U_\mu$  and  $(d - 1)$  transfer tensors  $B_t$ . Hence, if  $r = \max\{r_t : t \in \mathcal{T}\}$  and  $n = \max\{n_1, \dots, n_d\}$ , the storage requirements are bounded by

$$dnr + (d - 2)r^3 + r^2, \quad (9)$$

where we have used that the transfer tensor at the root can actually be considered as a matrix of size at most  $r \times r$ . While the complexity bound (9) appears to indicate linear dependence in  $d$ , it is important to note that such a conclusion only holds when the maximal hierarchical rank  $r$  can be assumed to stay constant as  $d$  increases. This is a rather strong assumption that is satisfied only in specific applications.

### 3. BASIC FUNCTIONALITY OF THE TOOLBOX

To conveniently work with tensors in HTD, we have implemented a new MATLAB class `htensor`, inspired by the classes `ktensor` (for tensors in CP decomposition) and `ttensor` (for tensors in Tucker decomposition) available in the Tensor Toolbox [Bader and Kolda 2006]. In the following, we describe the structure of `htensor` as well as its basic functionality.

#### 3.1. Fields and Properties of `htensor`

An instance of `htensor` contains two arrays specifying the dimension tree, an orthogonalization flag, as well as the matrices and transfer tensors representing an HTD corresponding to this dimension tree.



Each node of the dimension tree is associated with an index  $i \in \{1, \dots, 2d - 1\}$ , such that each child node has a larger index than the parent node. Consequently, the root node has index 1. The  $(2d - 1) \times 2$  integer array `children` specifies the structure of the dimension tree as follows: `children(i, 1)` is the left child of node  $i$ , and `children(i, 2)` is the right child of node  $i$ . Both entries are zero if node  $i$  is a leaf node. The  $1 \times d$  integer array `dim2ind` gives the index of the leaf node associated with each mode  $\mu = 1, \dots, d$ . The matrices  $U_t$  and transfer tensors  $B_t$  are stored in the cell arrays `U` and `B`, respectively. Note that `U{i}` is a matrix if  $i$  is a leaf node and an empty array for any other node. Finally, the boolean flag `is_orthog` indicates whether the HTD is orthogonalized; see Section 4.3.

**htensor for  $d = 4$  (see also Example 2.2):**

```
x.children: [2, 3; 4, 5; 6, 7; 0, 0; 0, 0; 0, 0; 0, 0]
x.dim2ind: [4 5 6 7]
x.U: {[ ] [ ] [ ] [4x4 double] [5x4 double] [6x6 double] [7x3 double]}
x.B: {[4x5 double] [4x4x4 double] [6x3x5 double] [ ] [ ] [ ] [ ]}
x.is_orthog: false
```

### 3.2. Constructors of `htensor`

There are several ways to construct an `htensor` instance. In the following, we only illustrate the most common ways and refer to the documentation, such as `help htensor/htensor`, for more details.

#### Examples for constructors of `htensor`:

```
x = htensor([4 5 6 7]) constructs a zero htensor of size  $4 \times 5 \times 6 \times 7$ .
x = htensor([4 5 6 7], 'TT') constructs a zero htensor of size  $4 \times 5 \times 6 \times 7$ , with
a degenerate, TT-like dimension tree.
x = htensor({U1, U2, U3}) constructs an htensor from the CP tensor defined by
 $\mathcal{X}(i_1, i_2, i_3) = \sum_{j=1}^R U_1(i_1, j) U_2(i_2, j) U_3(i_3, j)$ .
x = htenones([4 5 6 7]) constructs an htensor of size  $4 \times 5 \times 6 \times 7$ , with all
entries one.
x = htenrandn([4 5 6 7]) constructs an htensor of size  $4 \times 5 \times 6 \times 7$ , with random
ranks and random entries.
```

By default, any `htensor` has a balanced dimension tree. The motivation for the option 'TT' is to resemble the structure of the TTD (tensor train decomposition). However, it is important to note that there is no exact correspondence between HTD and TTD, as TTD does not require the storage of basis matrices. An arbitrary dimension tree can be generated by supplying the fields `children` and `dim2ind` to the constructor. Users of the Tensor Toolbox may also provide a `ktensor` for constructing an `htensor` from a CP decomposition, instead of providing the factors in a cell array as illustrated before.

### 3.3. Basic Functionality of `htensor`

Kressner and Tobler [2012, Table 1] contains all basic functions for working with `htensor` objects. The following example illustrates their use for a  $5 \times 4 \times 6 \times 3$  tensor  $\mathcal{X}$  in HTD as in Example 2.2.

`x(1, 3, 4, 2)` returns the entry  $\mathcal{X}_{1,3,4,2}$ .  
`x(1, 3, :, :)` returns an `htensor` representing the  $6 \times 3$  tensor of the slice  $\mathcal{X}_{1,3,:,\cdot}$ .  
`full(x)` returns the full tensor represented by  $\mathcal{X}$ .  
`x(:)` returns the vectorization of  $\mathcal{X}$ .  
`size(x)` returns the array of dimensions,  $[5, 4, 6, 3]$ .  
`ndims(x)` returns the order of the tensor, 4.  
`disp(htenrandn([5 4 6 3]))` returns the tree structure and the sizes of the transfer tensors/basis matrices.  
`disp_all(x)` additionally displays all transfer tensors and basis matrices.  
`spy(x)` displays the dimension tree with spy plots of  $U_t$  and  $B_t$ .  
`plot_sv(x)` displays the dimension tree with semi-log plots of the singular values of the matricizations at each node.

#### 4. BASIC OPERATIONS

This section describes algorithms and implementations for a range of typically required basic operations.

##### 4.1. $\mu$ -Mode Matrix Products

Given a tensor  $\mathcal{X} \in \mathbb{C}^{n_1 \times \dots \times n_d}$ , the  $\mu$ -mode product with a matrix  $A \in \mathbb{C}^{m \times n_\mu}$  is defined via the  $\mu$ -mode matricization.

$$\mathcal{Y} = A \circ_\mu \mathcal{X} \quad \Leftrightarrow \quad Y^{(\mu)} = AX^{(\mu)}$$

This operation can easily be performed if  $\mathcal{X}$  is in HTD: The  $\mu$ th basis matrix  $U_\mu$  is simply replaced by  $A U_\mu$ . We provide a function `ttm` for performing this operation. The calling sequence of `ttm` is nearly identical with the function `ttm` in the Tensor Toolbox. A notable difference is that we allow the matrix  $A$  to be provided implicitly as a handle to a function that returns the product of  $A$  with the input matrix. For example, `y = ttm(x, @(x)(fft(x)), 2)` applies the fast Fourier transformation to  $\mathcal{X}$  in mode 2.

##### 4.2. Addition

The addition of tensors in HTD can be performed at no arithmetic cost by a simple embedding. The underlying principle can easily be seen for two factorized matrices  $A = U_A \Sigma_A V_A^H$  and  $B = U_B \Sigma_B V_B^H$ :


$$A + B = \begin{bmatrix} U_A & U_B \end{bmatrix} \begin{bmatrix} \Sigma_A & 0 \\ 0 & \Sigma_B \end{bmatrix} \begin{bmatrix} V_A & V_B \end{bmatrix}^H.$$

This embedding is performed similarly for the addition of two or more tensors in HTD, by concatenation of the leaf matrices and a block diagonal embedding of the transfer tensors. We refrain from giving a technical description and refer to Figure 3 for an illustration. It is important to note that the storage requirements grow cubically in the number of tensors to be added if the block diagonal structure of the transfer tensors is not exploited. Such an alternative method is discussed in Section 6.3.

##### 4.3. Orthogonalization

An HTD of a tensor  $\mathcal{X}$  is called *orthogonalized* if the columns of  $U_t$  form an orthonormal basis for each node  $t$  except for the root node. Recall that the basis matrices  $U_t$  are defined recursively according to (6). As will be seen later, orthogonalized HTDs simplify some important operations related to tensor contraction. Moreover, they reduce the risk of numerical cancellation.



Fig. 3. Addition of four tensors  $\mathcal{X}_1 + \mathcal{X}_2 + \mathcal{X}_3 + \mathcal{X}_4$  in HTD.

We illustrate the process of orthogonalization for the tensor in standard HTD from Example 2.2:

$$\text{vec}(\mathcal{X}) = (U_4 \otimes U_3 \otimes U_2 \otimes U_1)(B_{34} \otimes B_{12})B_{1234}.$$

In the first step, QR decompositions of the basis matrices are performed:  $U_t = \tilde{U}_t R_t$  for  $t = 1, \dots, 4$ . Here and in the following, “economic” QR decompositions [Golub and Van Loan 1996] are performed, that is,  $U_t$  and  $\tilde{U}_t$  have the same number of columns. Propagating the factors  $R_t$  into the transfer matrices results in

$$\text{vec}(\mathcal{X}) = (\tilde{U}_4 \otimes \tilde{U}_3 \otimes \tilde{U}_2 \otimes \tilde{U}_1)(\hat{B}_{34} \otimes \hat{B}_{12})B_{1234}$$

with  $\hat{B}_{34} := (R_4 \otimes R_3)B_{34}$ ,  $\hat{B}_{12} := (R_2 \otimes R_1)B_{12}$ . In the next step, QR decompositions  $\hat{B}_{34} = \tilde{B}_{34}R_{34}$ ,  $\hat{B}_{12} = \tilde{B}_{12}R_{12}$  are performed, resulting in

$$\text{vec}(\mathcal{X}) = (\tilde{U}_4 \otimes \tilde{U}_3 \otimes \tilde{U}_2 \otimes \tilde{U}_1)(\tilde{B}_{34} \otimes \tilde{B}_{12})\tilde{B}_{1234} \quad (10)$$

with  $\tilde{B}_{1234} := (R_{34} \otimes R_{12})B_{1234}$ . Clearly, (10) constitutes an orthogonalized HTD, completing the orthogonalization procedure.

This procedure easily extends to the general case; see Kressner and Tobler [2012, Algorithm 1]. Properly implemented, orthogonalization requires  $\mathcal{O}(dnr^2 + dr^4)$  operations. Unless  $r$  is very small, the factor  $dr^4$ , caused by the QR decompositions of the transfer matrices, will be dominant.

In the htucker toolbox, orthogonalization is performed by calling `x = orthog(x)`. On return, the flag `is_orthog` of the htensor object `x` is set to true. This prevents unnecessary orthogonalization in subsequent calls to `orthog`. When this property is destroyed by an operation, such as addition and  $\mu$ -mode matrix products, the flag `is_orthog` is set to false.

## 5. TENSOR-TENSOR CONTRACTION

The contraction of two tensors  $\mathcal{X} \in \mathbb{C}^{m_1 \times \dots \times m_c}$  and  $\mathcal{Y} \in \mathbb{C}^{n_1 \times \dots \times n_d}$  is a fundamental operation which generalizes the concepts of inner product, outer product, and matrix multiplication. In its most general form, we select a tuple of modes  $s = (i_1, \dots, i_p)$  from  $\mathcal{X}$  and another tuple of modes  $t = (j_1, \dots, j_p)$  from  $\mathcal{Y}$ . Then the corresponding

contraction of  $\mathcal{X}$  and  $\mathcal{Y}$  amounts to taking the inner product with respect to each pair  $(i_\mu, j_\mu)$  of selected modes. This implicitly assumes that the sizes of the selected modes match, that is,  $n_{i_\mu} = m_{j_\mu}$  for  $\mu = 1, \dots, p$ . Contraction results in a tensor  $\mathcal{Z}$  whose order equals the number of nonselected modes. For example, for  $c = 4, d = 3$ , and  $s = (3, 1), t = (2, 3)$ , the contracted tensor  $\mathcal{Z} \in \mathbb{C}^{m_2 \times m_4 \times n_1}$  is given by

$$\mathcal{Z}_{i_1, i_2, i_3} = \langle \mathcal{X}, \mathcal{Y} \rangle_{(3,1),(2,3)} := \sum_{j=1}^{n_3} \sum_{k=1}^{n_1} \bar{\mathcal{X}}_{k, i_1, j, i_2} \mathcal{Y}_{i_3, j, k}. \quad (11)$$

A matrix-matrix multiplication  $X^H Y$  for  $X \in \mathbb{C}^{n_1 \times n_2}$ ,  $Y \in \mathbb{C}^{m_1 \times m_2}$  with  $n_1 = m_1$  can be seen as the contraction corresponding to  $s = (1), t = (1)$ . Conversely, a general tensor-tensor contraction of  $\mathcal{X} \in \mathbb{C}^{m_1 \times \dots \times m_c}$  and  $\mathcal{Y} \in \mathbb{C}^{n_1 \times \dots \times n_d}$  along  $p$  selected pairs of nodes can be defined via matrix-matrix multiplication.


$$Z^{(\bar{t})} = \left( \langle \mathcal{X}, \mathcal{Y} \rangle_{(t;s)} \right)^{(\bar{t})} := (X^{(t)})^H Y^{(s)}, \quad \text{with } \bar{t} = (1, 2, \dots, c-p) \quad (12)$$

Note that the matricization needs to respect the order of the modes in  $t$  and  $s$ ; see also Section 2.1.

### 5.1. Tensor Network Diagrams

In the following, we briefly introduce tensor network diagrams (also called Penrose diagrams) to conveniently describe algorithms for contractions of tensors in HTD; see also Holtz et al. [2012] and Huckle et al. [2013]. Such a diagram represents a tensor in terms of contractions of other tensors. Each node in the diagram represents a tensor and each edge represents a mode. An edge connecting two nodes corresponds to the contraction of these tensors in the associated pair of modes. We allow for dangling edges, which are connected to only one node and correspond to modes that are not contracted. Hence, the order of the tensor represented by the network is given by the number of dangling edges.

To illustrate the concept of tensor network diagrams, let us consider the contraction given in (11). This is represented by the following diagram.



The node  $\mathcal{X}$  has four edges, each corresponding to one of the four different modes of the tensor  $\mathcal{X}$ . Similarly, the node  $\mathcal{Y}$  has three edges. Since mode 3 of  $\mathcal{X}$  is contracted with mode 2 of  $\mathcal{Y}$ , the nodes share the corresponding edge. Similarly, the contraction of modes 1 and 3 is represented by another shared edge.

Some further examples of tensor network diagrams are given in Figure 4. Note that we will only indicate the precise mode(s) belonging to an edge when necessary. In the case of HTD, each edge connecting two nodes corresponds to a matricization  $X^{(t)}$  for some  $t \in \mathcal{T}$ .

### 5.2. Inner Product and Norm for Tensors in HTD

The inner product of two tensors  $\mathcal{X}, \mathcal{Y} \in \mathbb{C}^{n_1 \times \dots \times n_d}$  is an important special case of contraction. We have

$$\langle \mathcal{X}, \mathcal{Y} \rangle = \sum_{i_1=1}^{n_1} \dots \sum_{i_d=1}^{n_d} \bar{\mathcal{X}}_{i_1, \dots, i_d} \mathcal{Y}_{i_1, \dots, i_d}$$




Fig. 4. Tensor network diagrams representing: (i) a vector; (ii) a matrix; (iii) a matrix-matrix multiplication; (iv) a tensor in Tucker decomposition; (v) a tensor in HTD.




Fig. 5. Inner product of two tensors of order 4 in HTD.

or, equivalently,  $\langle \mathcal{X}, \mathcal{Y} \rangle = \langle \text{vec}(\mathcal{X}), \text{vec}(\mathcal{Y}) \rangle$ . In terms of tensor network diagrams, this operation corresponds to a pairwise connection of the dangling edges of  $\vec{\mathcal{X}}$  and  $\vec{\mathcal{Y}}$ .

To illustrate how to evaluate inner products of tensors in HTD efficiently, we first consider two tensors of order 4.

$$\langle \mathcal{X}, \mathcal{Y} \rangle = (B_{1234}^x)^H (B_{34}^x \otimes B_{12}^x)^H (U_4^x \otimes U_3^x \otimes U_2^x \otimes U_1^x)^H (U_4^y \otimes U_3^y \otimes U_2^y \otimes U_1^y) (B_{34}^y \otimes B_{12}^y) B_{1234}^y$$

This product is evaluated from inside to outside or, when considering the hierarchical tree, from the leaves to the root node. In a first step, the matrices

$$M_t = (U_t^x)^H U_t^y, \quad t = 1, \dots, 4,$$

are computed. Then the product

$$M_t = (U_t^x)^H U_t^y = (B_t^x)^H (M_{t_r} \otimes M_{t_l}) B_t^y \quad (13)$$

is computed for  $t = \{1, 2\}$  and  $t = \{3, 4\}$ . Note that the matrix  $M_{t_r} \otimes M_{t_l}$  is not formed explicitly but the matrices  $M_{t_l}$  and  $M_{t_r}$  are applied to the modes 1 and 2 of  $B_t^y$ , respectively. The product (13) then requires  $6r^4$  operations if both  $\mathcal{X}$  and  $\mathcal{Y}$  have constant hierarchical rank  $r$ . In the last step,  $\langle \mathcal{X}, \mathcal{Y} \rangle$  is obtained by evaluating (13) for  $t = \{1, 2, 3, 4\}$ . Figure 5 illustrates the described procedure with tensor network diagrams.

The generalization to tensors of arbitrary order is straightforward; see Kressner and Tobler [2012, Algorithm 2]. It is implemented in the `htucker` function `innerprod`. In total, forming the inner product of two tensors with constant hierarchical rank  $r$  requires  $6(d-1)r^4 + \sum_{\mu=1}^d 2n_{\mu}r^2$  operations.

In principle, the Euclidean norm of a tensor  $\mathcal{X}$  in HTD can be calculated from  $\|\mathcal{X}\|_2 = \sqrt{\langle \mathcal{X}, \mathcal{X} \rangle}$  by means of an inner product. However, it is well known that such an approach suffers numerical instabilities and may introduce an error proportional to the square root of machine precision. A usually more accurate alternative is to first orthogonalize the HTD of  $\mathcal{X}$  and then compute the norm.

### 5.3. General Contraction of Tensors in HTD

In tensor network diagrams, a general contraction of two tensors in HTD is performed by connecting the corresponding pairs of dangling edges. This will create a tensor network with cycles which need to be eliminated. This elimination is performed by successive contraction of tensors, similarly as earlier, until the network becomes a tree. Note that this procedure can only be organized efficiently if the maximum degree of all intermediate tensor networks does not become too large. In turn, this imposes certain restrictions on the input data to the function `ttt` that implements a general contraction. For details, we refer to Tobler [2012, Section 3.5.2].

### 5.4. Reduced Gramians of a Tensor in HTD

An important application of contractions is the calculation of reduced Gramians, which are defined as follows. For every  $t \in \mathcal{T}$ , the matrix  $U_t$  defined in (6) contains a basis for the column span of the matricization  $X^{(t)}$ . Hence, there is a matrix  $V_t$  such that  $X^{(t)} = U_t V_t^H$ . The *reduced Gramian* at  $t$  is then defined as the Hermitian positive semi-definite matrix  $G_t = V_t^H V_t \in \mathbb{C}^{r_t \times r_t}$ .

Reduced Gramians are a central tool in the truncation of tensors. For example, they provide an efficient way to compute the singular values of  $X^{(t)}$ , which is used in the function `plot_sv`. From

$$X^{(t)}(X^{(t)})^H = U_t V_t^H V_t U_t^H = U_t G_t U_t^H, \quad (14)$$

it follows that the singular values of  $X^{(t)}$  are the square roots of the eigenvalues of the reduced Gramian  $G_t$ , provided that  $U_t^H U_t = I_{r_t}$ . This condition is always satisfied after the HTD has been orthogonalized.

In the following, we briefly discuss the computation of a reduced Gramian  $G_t$  via contraction. The standard, unreduced Gramian corresponds to the contracted product of  $\mathcal{X}$  with itself along the modes  $t^c = \{1, \dots, d\} \setminus t$ :  $\langle \mathcal{X}, \mathcal{X} \rangle_{t^c}$ , for which the matricization is given by (14); see also Figure 6 for an illustration. It can be seen from the figure that Gramians are closely related to contraction. In particular, this relation implies that  $G_t$  can be calculated by a sequence of matrix-tensor and tensor-tensor products.

Typically, not the reduced Gramian at one node  $t$  is required, but reduced Gramians  $G_t$  for all nodes  $t \in \mathcal{T}$ . These can be computed simultaneously by exploiting relations between different reduced Gramians. The relationship between the reduced Gramians  $G_t$  and  $G_{t_r}$ , where  $t_r$  is the right child node of  $t$ , is illustrated in Figure 7 for the case of a general non-leaf node  $t$ . From the tensor network diagram, it can be seen that

$$\begin{aligned} G_{t_r} &= (B_t^{((1,3))})^H (G_t \otimes U_{t_l}^H U_{t_l}) B_t^{((1,3))}, \\ G_{t_l} &= (B_t^{((2,3))})^H (G_t \otimes U_{t_r}^H U_{t_r}) B_t^{((2,3))}. \end{aligned} \quad (15)$$

Formally setting  $G_{t_{\text{root}}} = 1$ , this defines a recursive algorithm for the calculation of all reduced Gramians. An efficient algorithm for calculating  $M_t = U_t^H U_t$  is already




Fig. 6. Reduced Gramian of a tensor of order 8. The reduced Gramian  $G_t$  corresponds to the left subnetwork encircled by a dashed line.

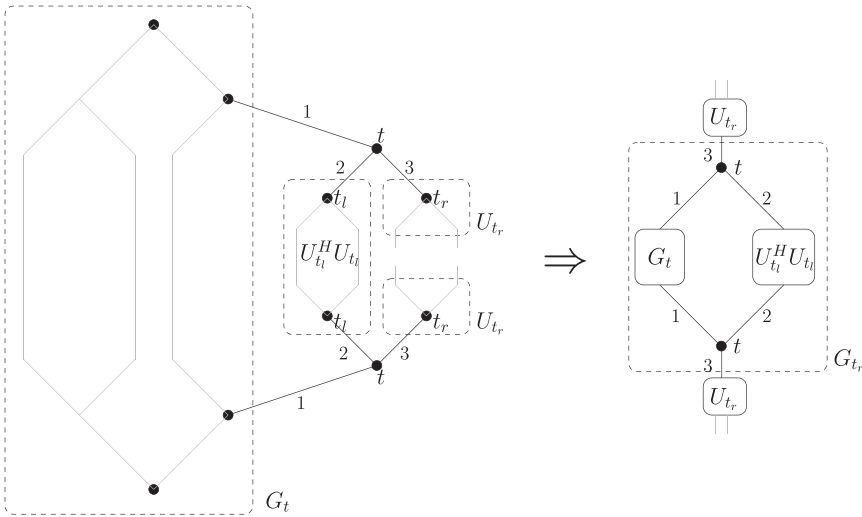


Fig. 7. How to calculate  $G_{tr}$  from  $G_t$ ,  $U_{t_l}$  and  $B_t$ , where  $t$  is a general non-leaf node. The gray lines represent arbitrary subtrees.

described in Section 5.2. Note that all  $M_t$  are identity matrices for a tensor in orthogonalized HTD.

For a general HTD, the computation of the reduced Gramians requires  $O(dnr^2 + (d - 2)r^4)$  operations. For an orthogonalized HTD, this reduces to  $O((d - 2)r^4)$  operations (but orthogonalization itself requires  $O(dnr^2 + (d - 2)r^4)$  operations).

$G = \text{gramians}(x)$ ; Reduced Gramians of  $\mathcal{X}$  in cell array  $G$ , orthogonalizing the HTD of  $\mathcal{X}$  if necessary.

$G = \text{gramians\_nonorthog}(x)$ ; Reduced Gramians of  $\mathcal{X}$  in cell array  $G$ , without orthogonalizing.

$sv = \text{singular\_values}(x)$ ; Singular values of  $\mathcal{X}$  in cell array  $sv$ .

## 6. TRUNCATION OF TENSORS

Truncation of tensors to HTD is one of the most important and most frequently used operations in `htucker`.

### 6.1. Truncation of Explicit Tensors

We start with the truncation of an explicitly given tensor  $\mathcal{X} \in \mathbb{C}^{n_1 \times \dots \times n_d}$ . Although this situation is limited to small dimensions/sizes, it provides a gentle introduction and illustration of the general concepts. Truncation to HTD is done by successive projections to the subspaces  $\text{span}(W_t)$ , which typically represent approximations to the column spaces of  $X^{(t)} \in \mathbb{C}^{n_t \times n_{t^c}}$ . For a subset  $t \subset \{1, \dots, d\}$ , we define  $n_t := \prod_{\mu \in t} n_\mu$ . We require  $W_t \in \mathbb{C}^{n_t \times r_t}$  to have orthonormal columns and define the orthogonal projections

$$\pi_t := W_t W_t^H. \quad (16)$$

In the following, we use the shorthand notation  $\pi_t \mathcal{X}$  for  $\pi_t \circ_t \mathcal{X}$ . As shown in Grasedyck [2010b, Lemma 3.15], applying these projections in the correct order leads to a tensor in HTD, with hierarchical ranks bounded by  $r_t$ .

**6.1.1. Root-to-Leaves Truncation.** The simplest way to construct the projections  $\pi_t$  in (16) is to let each matrix  $W_t$  contain the  $r_t$  dominant left singular vectors of the corresponding matricization  $X^{(t)}$ . To obtain a tensor in HTD,  $\tilde{\mathcal{X}} \in \mathcal{H}\text{-Tucker}((r_t)_{t \in \mathcal{T}})$ , the projections need to be applied from the root node to the leaves. The computation for the general case is described in Grasedyck [2010b, Algorithm 1] and Kressner and Tobler [2012, Algorithm 4]. Note that the HTD of the resulting tensor is not orthogonalized, only the matrices in the leaf nodes have orthonormal columns. Setting  $n = \max_t n_t$ , the computational complexity of root-to-leaves truncation is of order  $dn^{3d/2}$  in the case of a balanced tree. The resulting tensor  $\tilde{\mathcal{X}}$  satisfies the following error bound [Grasedyck 2010b, Theorem 3.11]

$$\|\mathcal{X} - \tilde{\mathcal{X}}\|_2 \leq \sqrt{\sum_{t \in \mathcal{T}'} \delta_{r_t}(X^{(t)})^2} \leq \sqrt{2d-3} \|\mathcal{X} - \mathcal{X}_{\text{best}}\|_2, \quad (17)$$

where  $\mathcal{X}_{\text{best}}$  represents the best approximation of  $\mathcal{X}$  in  $\mathcal{H}\text{-Tucker}((r_t)_{t \in \mathcal{T}})$ ,  $\mathcal{T}' := \mathcal{T} \setminus \{t_{\text{root}}, t_{\text{child}}\}$  where  $t_{\text{child}}$  is a child of the root node  $t_{\text{root}}$ , and

$$\delta_{r_t}(X^{(t)})^2 := \sum_{j=r_t+1}^{n_t} \sigma_j(X^{(t)})^2. \quad (18)$$

**Remark 6.1.** The error bound (17) allows us to choose the hierarchical ranks  $(r_t)_{t \in \mathcal{T}}$  such that a certain error bound is satisfied

$$\begin{aligned} \|\mathcal{X} - \tilde{\mathcal{X}}\|_2 &\leq \epsilon_{\text{abs}}, \quad \text{choose } r_t \text{ such that } \delta_{r_t}(X^{(t)}) \leq \frac{\epsilon_{\text{abs}}}{\sqrt{2d-3}} \quad \forall t \in \mathcal{T} \setminus \{t_{\text{root}}\}, \\ \|\mathcal{X} - \tilde{\mathcal{X}}\|_2 &\leq \epsilon_{\text{rel}} \|\mathcal{X}\|_2, \quad \text{choose } r_t \text{ such that } \delta_{r_t}(X^{(t)}) \leq \frac{\epsilon_{\text{rel}} \|\mathcal{X}\|_2}{\sqrt{2d-3}} \quad \forall t \in \mathcal{T} \setminus \{t_{\text{root}}\}. \end{aligned}$$



Similar adaptive choices of the hierarchical ranks are possible for all other truncation methods discussed in the following.

```

opts.max_rank = 10; maximal rank at truncation, mandatory argument.
opts.rel_eps = 1e-6; maximal relative truncation error, optional argument.
opts.abs_eps = 1e-6; maximal absolute truncation error, optional argument.
Condition max_rank takes precedence over rel_eps and abs_eps.
y = htensor.truncate_rtl(x, opts); takes a MATLAB multidimensional array
and returns the truncation to lower rank HTD.

```

**6.1.2. Leaves-to-Root Truncation.** Root-to-leaves truncation is very costly, the most expensive part being the computation of the singular value decomposition of every  $X^{(t)} \in \mathbb{C}^{n_t \times n_{tc}}$ , where both  $n_t$  and  $n_{tc}$  can become very large. Leaves-to-root truncation can be considerably faster. It is described in Grasedyck [2010b, Algorithm 2] and Kressner and Tobler [2012, Algorithm 5], and implemented in the function `x = htensor.truncate_ltr(x, opts)` of the toolbox. The computational complexity of leaves-to-root truncation is  $O(dn^{d+1})$ , which is a significant reduction compared to the root-to-leaves method, while the error bound (17) still holds; see Lemma B.2 in Kressner and Tobler [2012]. Moreover, the resulting tensor  $\tilde{\mathcal{X}}$  is in orthogonalized HTD.

## 6.2. Truncation of $\mathcal{H}$ -Tucker Decomposition to Lower Rank

The truncation of a tensor which is already given in HTD to a tensor in HTD of lower rank is an essential operation in most algorithms based on this format. In Section 6.2.1, we describe an efficient method for performing such a truncation. This will be the method of choice for general tensors. However, for structured tensors resulting, for example, from the addition of several tensors in HTD, a different approach described in Section 6.2.2 is preferable.

**6.2.1. Truncation of a Tensor in HTD.** Truncation of a tensor  $\mathcal{X}$  in HTD can be performed by a fairly straightforward adaptation of the root-to-leaves method. For this purpose, we recall that Section 5.4 describes an efficient method for computing the reduced Gramians  $G_t$  in the decomposition

$$X^{(t)}(X^{(t)})^H = U_t G_t U_t^H,$$

where  $U_t$  has orthonormal columns and is implicitly represented as illustrated in Figure 7. After orthogonalizing the HTD of  $\mathcal{X} \in \mathcal{H}\text{-Tucker}((k_t)_{t \in \mathcal{T}})$  and calculating the reduced Gramians, we compute an orthonormal basis  $S_t \in \mathbb{C}^{k_t \times r_t}$  for the  $r_t$  dominant eigenvectors of the symmetric matrix  $G_t$ . As before, we define  $W_t := U_t S_t$  and obtain the truncated tensor  $\tilde{\mathcal{X}}$  from subsequent application of the projections  $\pi_t = W_t W_t^H$ .

To illustrate how these projections can be applied to a tensor in HTD, let us consider the example of a tensor  $\mathcal{X}$  of order 4.

$$\begin{aligned}
\text{vec}(\tilde{\mathcal{X}}) &= (W_4 W_4^H \otimes W_3 W_3^H \otimes W_2 W_2^H \otimes W_1 W_1^H) (W_{34} W_{34}^H \otimes W_{12} W_{12}^H) \text{vec}(\mathcal{X}) \\
&= (U_4 S_4 \otimes U_3 S_3 \otimes U_2 S_2 \otimes U_1 S_1) \cdots \\
&\quad \cdots \underbrace{((S_4^H \otimes S_3^H) B_{34} S_{34})}_{=:\tilde{B}_{34}} \underbrace{(S_2^H \otimes S_1^H) B_{12} S_{12}}_{=:\tilde{B}_{12}} \underbrace{(S_{34}^H \otimes S_{12}^H) B_{1234}}_{=:\tilde{B}_{1234}}
\end{aligned}$$

Hence, an HTD for  $\tilde{\mathcal{X}}$  is obtained by updating the leaf matrices  $\tilde{U}_1 := U_1 S_1, \dots, \tilde{U}_4 := U_4 S_4$ , and the transfer matrices. Note that the matrices  $W_t$  are never calculated explicitly.

Algorithm 6 in Kressner and Tobler [2012], which describes this procedure, is implemented in `truncate_std(x, opts)`, and requires  $O(dnr^2 + dr^4)$  operations. As this algorithm is mathematically identical to the explicit root-to-leaves algorithm described in Section 6.1.1, the error bound (17) holds. Note that the resulting tensor  $\tilde{\mathcal{X}}$  is not in orthogonalized HTD.

**6.2.2. Truncation of a Tensor in HTD without Initial Orthogonalization.** The method of truncation introduced in the last section represents the default method for truncating a tensor in HTD. However, in certain situations, it can be beneficial to exploit additional structure in the HTD. For example, a tensor resulting from addition of tensors in HTD has block diagonal transfer tensors. In the standard method of truncation, such structures are immediately destroyed by the initial orthogonalization step. In the following, we discuss a method that avoids this step.

In a first step, the reduced Gramians  $G_t$  in the decomposition

$$X^{(t)}(X^{(t)})^H = U_t G_t U_t^H$$

are calculated without the initial orthogonalization. Note, however, that the singular value decomposition of  $X^{(t)}$  cannot be computed directly from  $G_t$ , as the columns of  $U_t$  are not orthonormal.

In a second step, the proposed method successively orthonormalizes the matrices  $U_t$ . Let us first consider the leaf nodes  $t$ , for which we compute the  $r_t$  dominant left singular vectors of  $X^{(t)}$  as follows: Compute the QR decomposition  $U_t =: Q_t R_t$ , and determine the matrix  $S_t$  containing the  $r_t$  dominant eigenvectors of  $R_t G_t R_t^H$ . Then the projection  $\pi_t = W_t W_t^H$ , with  $W_t = Q_t S_t$ , is applied to  $\mathcal{X}$ . Note that the updated leaf nodes  $\tilde{U}_t := Q_t S_t$  are orthonormal.

Non-leaf nodes are processed in a similar manner with a recursive algorithm, traversing the tree such that every parent node is visited after its child nodes. Assume we are at node  $t$ , and let  $\tilde{U}_t$  account for all updates from previous operations on the descendants of  $t$ . Based on the original decomposition  $X^{(t)} = U_t V_t^H$ , we set  $\tilde{X}_t^{(t)} := \tilde{U}_t V_t^H$  and observe that the corresponding Gramian takes the form

$$\tilde{X}_t^{(t)}(\tilde{X}_t^{(t)})^H = \tilde{U}_t G_t \tilde{U}_t^H.$$

To orthogonalize  $\tilde{U}_t = (\tilde{U}_{t_r} \otimes \tilde{U}_{t_l}) \tilde{B}_t$ , it is sufficient to calculate the QR decomposition of  $\tilde{B}_t = (S_{t_r}^H R_{t_r} \otimes S_{t_l}^H R_{t_l}) B_t$ , as the columns of  $\tilde{U}_{t_l}$  and  $\tilde{U}_{t_r}$  are orthonormal. Then, we calculate the  $r_t$  dominant left singular vectors of  $\tilde{X}_t^{(t)}$  as in the case of the leaf nodes.

A more detailed description of truncation to HTD without initial orthogonalization can be found in Kressner and Tobler [2012, Algorithm 7] and is implemented in `truncate_nonorthog(x, opts)` in the toolbox. The result of this algorithm satisfies practically the same error bound as in (17); see also Lemma B.3 in Kressner and Tobler [2012].

### 6.3. Combined Addition and Truncation

As explained in Section 4.2, the addition of tensors in HTD leads to a significant growth of the hierarchical ranks. For example, the sum of  $s$  tensors of hierarchical ranks  $r$  has hierarchical ranks  $sr$ . Truncation of this tensor back to hierarchical rank  $r$  requires  $O(dns^2r^2 + ds^4r^4)$  operations, which is too expensive unless  $s$  is very small.




Fig. 8. Execution times for truncating a sum of tensors in HTD. Red: Standard truncation of the sum. Green: Truncation of the sum with using combined addition and truncation, as described in Section 6.3. Blue: Successive addition and truncation; see (19).

A less expensive alternative is to add the  $s$  tensors successively and truncate immediately after each addition. After setting  $\tilde{\mathcal{Y}}_1 := \mathcal{X}_1$ , we compute for  $j = 1, \dots, s-1$ .

$$\begin{aligned} &\text{Form } \mathcal{Y}_{j+1} := \tilde{\mathcal{Y}}_j + \mathcal{X}_j \\ &\text{Truncate } \mathcal{Y}_{j+1} \text{ to } \tilde{\mathcal{Y}}_{j+1} \end{aligned} \quad (19)$$

However, one can easily construct examples for which this scheme suffers severe cancellation (see `example_cancellation.m` in the toolbox).

To avoid cancellation and still increase efficiency, we propose to apply truncation without initial orthogonalization (Section 6.2.2) directly to the sum of tensors and exploit the block diagonal structures illustrated in Figure 3. This results in significant savings when calculating the reduced Gramians. Computing the reduced Gramians  $G_t$  at a non-leaf node requires only  $O(s^2 r^4)$  instead of  $O(s^4 r^4)$  operations. Hence, the computational cost of the whole addition and truncation process reduces to  $O(dns^2 r^2 + ds^2 r^4 + ds^3 r^3)$ .

With a numerical experiment we examine the execution time required for the addition and truncation of  $s$  random tensors of order  $d = 5$ , with size  $n = 500$  and rank  $r = 20$ . The number of summands  $s$  varies between 2 and 10 (see Figure 8). This numerical experiment was performed in MATLAB, version 7.6, on an Intel Xeon CPU E31225 with 3 GHz and 6MB L2 Cache and 8GB RAM. The execution time of the new method increases proportionally with  $s^2$ , indicating that the term  $s^3 r^3$  does not dominate the cost for this rather typical setting. Note that the execution time of the new method is relatively high for small  $s$ . However, this only reflects the additional overhead of this method in a MATLAB implementation; even for  $s = 2$ , the operation count of the new method is smaller compared to applying standard truncation to the sum of tensors.

```
xt = truncate_sum({x1, x2, x3}, opts);
```

takes htensor objects  $\mathcal{X}_1, \mathcal{X}_2, \mathcal{X}_3$  and returns a truncated htensor  $\tilde{\mathcal{X}} \approx \mathcal{X}_1 + \mathcal{X}_2 + \mathcal{X}_3$ .

## 7. ELEMENTWISE MULTIPLICATION

The elementwise multiplication of two tensors is an important operation in connection with function-related tensors and can be performed efficiently for tensors in HTD. In

particular, this operation corresponds to the pointwise multiplication  $f(\xi_1, \dots, \xi_d) = f_1(\xi_1, \dots, \xi_d)f_2(\xi_1, \dots, \xi_d)$  of the functions  $f_1, f_2$  represented by the two tensors.

For illustration, we first consider the elementwise multiplication of two low-rank matrices  $X = U^x S^x (V^x)^H$  and  $Y = U^y S^y (V^y)^H$ , with  $S^x, S^y \in \mathbb{C}^{r \times r}$ . Then the elementwise product (also called Hadamard product) can be written as

$$\begin{aligned} (X \star Y)_{ij} &:= X_{ij} Y_{ij} = \sum_{\alpha, \beta, \gamma, \delta} U_{i, \alpha}^x U_{i, \gamma}^y S_{\alpha, \beta}^x S_{\gamma, \delta}^y V_{j, \beta}^x V_{j, \delta}^y \\ &= (U^x \odot^T U^y) (S^x \otimes S^y) (V^x \odot^T V^y)^H, \end{aligned}$$

where  $\odot^T$  denotes a transposed variant of the Khatri-Rao product [Kolda and Bader 2009]. More specifically, for a matrix  $A \in \mathbb{C}^{n \times k}$  with rows  $a_j^T$  and a matrix  $B \in \mathbb{C}^{n \times r}$  with rows  $b_j^T$ , we define

$$A \odot^T B = \begin{bmatrix} a_1^T \\ a_2^T \\ \vdots \\ a_n^T \end{bmatrix} \odot^T \begin{bmatrix} b_1^T \\ b_2^T \\ \vdots \\ b_n^T \end{bmatrix} := \begin{bmatrix} a_1^T \otimes b_1^T \\ a_2^T \otimes b_2^T \\ \vdots \\ a_n^T \otimes b_n^T \end{bmatrix} \in \mathbb{C}^{n \times kr}.$$

Note that  $A \odot^T B = (A^T \odot B^T)^T$ , where  $\odot$  is the usual Khatri-Rao product.

For the elementwise multiplication of two tensors  $\mathcal{X}, \mathcal{Y}$  in HTD, the same technique can be used to construct the leaf matrices  $U_t$  and transfer tensors  $\mathcal{B}_t$  for an HTD of  $\mathcal{X} \star \mathcal{Y}$

$$U_t = U_t^x \odot^T U_t^y \text{ for leaf nodes } t \text{ and } \mathcal{B}_t = \mathcal{B}_t^x \otimes \mathcal{B}_t^y \text{ for non-leaf nodes } t,$$

where  $\otimes$  represents a direct generalization of the Kronecker product to tensors: For two tensors  $\mathcal{X} \in \mathbb{C}^{n_1 \times \dots \times n_d}, \hat{\mathcal{X}} \in \mathbb{C}^{\hat{n}_1 \times \dots \times \hat{n}_d}$  we define  $(\mathcal{X} \otimes \hat{\mathcal{X}})_{J_1, \dots, J_d} := \mathcal{X}_{i_1, \dots, i_d} \hat{\mathcal{X}}_{j_1, \dots, j_d}$ , with  $J_\mu = (i_\mu - 1)\hat{n}_\mu + j_\mu$ . As a consequence, the hierarchical rank  $r_t$  of  $\mathcal{X} \star \mathcal{Y}$  is the product  $r_t^x r_t^y$  of the hierarchical ranks of  $\mathcal{X}$  and  $\mathcal{Y}$ .

It would be useful to avoid the rank growth of  $\mathcal{X} \star \mathcal{Y}$  and directly calculate a truncated version, similarly as for the sum of  $s$  tensors in Section 6.3. Unfortunately, it is not clear how to transfer the ideas from Section 6.3 to the elementwise product; the Kronecker structure of the transfer tensors does not lead in an obvious way to a reduction of computational or storage cost. However, we can exploit the fact that the elementwise product is contained in the Kronecker product. More specifically, there is a  $(0, 1)$ -matrix  $J_n \in \mathbb{R}^{n^2 \times n}$  with orthonormal columns such that

$$a \star b = J_n^H (a \otimes b),$$

for any two vectors  $a, b \in \mathbb{C}^n$ . This extends in a direct fashion to tensors.

$$J^H(\mathcal{X} \otimes \mathcal{Y}) := (J_{n_d}^H \otimes \dots \otimes J_{n_1}^H)(\mathcal{X} \otimes \mathcal{Y}) = \mathcal{X} \star \mathcal{Y} \quad (20)$$

Hence, we can implicitly form the Kronecker product  $\mathcal{X} \otimes \mathcal{Y}$  in HTD and extract the elementwise product after truncation.

The HTD of the Kronecker product  $\mathcal{X} \otimes \mathcal{Y}$  of two tensors  $\mathcal{X}, \mathcal{Y}$  in HTD is particularly simple:

$$\begin{aligned} U_t &:= U_t^x \otimes U_t^y \quad \forall t \in \mathcal{L}(\mathcal{T}), \\ B_t &:= B_t^x \otimes B_t^y \quad \forall t \in \mathcal{N}(\mathcal{T}). \end{aligned}$$

This implies that the reduced Gramians have Kronecker structure,  $G_t = G_t^x \otimes G_t^y$ , as well as their singular value decompositions used in the standard truncation to lower rank. Consequently, this allows for a particularly efficient HTD truncation  $\mathcal{Z}$  of  $\mathcal{X} \otimes \mathcal{Y}$ . Using (20), the extracted tensor  $J^H \mathcal{Z}$  represents an approximation of  $\mathcal{X} \star \mathcal{Y}$  satisfying the error bound

$$\|\mathcal{X} \star \mathcal{Y} - J^H \mathcal{Z}\|_2 = \|J^H(\mathcal{X} \otimes \mathcal{Y} - \mathcal{Z})\|_2 \leq \|\mathcal{X} \otimes \mathcal{Y} - \mathcal{Z}\|_2 \leq \epsilon_{abs}.$$

Although the hierarchical ranks of  $J^H \mathcal{Z}$  are typically much smaller compared to  $\mathcal{X} \star \mathcal{Y}$ , the error bound given before is far from sharp. It is therefore recommended to truncate  $J^H \mathcal{Z}$  again after the extraction.

`z = x .* y` elementwise product of  $\mathcal{X}$  and  $\mathcal{Y}$ .  
`z = elem_mult(x, y, opts)` approximate elementwise product, with `opts` defined as in `truncate`.

## 8. HTD OF LINEAR OPERATORS ON TENSORS

The use of tensors in PDE-related applications often requires the efficient storage and application of a linear operator to tensors. In many cases, such a linear operator can be written as a short sum of Kronecker products:

$$\text{vec}(\mathcal{A}(\mathcal{X})) = \sum_{j=1}^R \left( A_j^{(d)} \otimes \cdots \otimes A_j^{(1)} \right) \text{vec}(\mathcal{X}), \quad \text{with } A_j^{(\mu)} \in \mathbb{C}^{m_\mu \times n_\mu}. \quad (21)$$

For example, a discretized Laplace operator in  $d$  dimensions takes this form with  $R = d$ ; see Example 8.1 that is given shortly. The operation (21) can be implemented by first applying  $\mu$ -mode matrix products (`ttm`) and then using an algorithm for computing a sum of tensors in HTD; see Section 6.3. In general, this is a reasonable approach. However, for particular linear operators, a much more efficient scheme can be devised.

This scheme is based on interpreting a linear operator as a tensor, an idea which goes back to the computational physics community [Schollwöck 2011]. For example, the operator in (21) can be vectorized into

$$\tilde{\mathcal{A}} = \sum_{j=1}^R \text{vec} \left( A_j^{(d)} \right) \otimes \cdots \otimes \text{vec} \left( A_j^{(1)} \right).$$

Note that this format is a CP decomposition. More generally, there is an isomorphism

$$\Psi : \mathfrak{L}(\mathbb{C}^{n_1 \times \cdots \times n_d}, \mathbb{C}^{m_1 \times \cdots \times m_d}) \rightarrow \mathbb{C}^{n_1 m_1 \times \cdots \times n_d m_d},$$

which takes the matrix representation  $\mathcal{A}_M \in \mathbb{C}^{(m_1 \cdots m_d) \times (n_1 \cdots n_d)}$  of a linear operator  $\mathcal{A}$ , and permutes and reshapes its entries into a tensor  $\tilde{\mathcal{A}} = \Psi(\mathcal{A})$  of order  $d$ .

The tensor  $\tilde{\mathcal{A}} = \Psi(\mathcal{A})$  can now be approximated in HTD by the methods described in this article. When applying a linear operator implicitly represented as a tensor in HTD with leaf bases  $U_t^A \in \mathbb{C}^{m_t n_t \times s_t}$  and transfer tensors  $B_t^A$ , it is convenient to reinterpret the columns of the leaf bases as matrices  $A_t^{(j)}$ :

$$U_t^A(:, j) = \text{vec} \left( A_t^{(j)} \right).$$

Then the application of  $\mathcal{A}$  to a tensor  $\mathcal{X}$  of conforming size in HTD with hierarchical ranks  $r_t$  again results in an HTD with

$$U_t = \left[ A_t^{(1)} U_t^x, \dots, A_t^{(s_t)} U_t^x \right], \quad B_t = B_t^A \otimes B_t^x.$$

Hence, the hierarchical ranks grow to  $s_t r_t$ , which illustrates the importance of keeping the hierarchical ranks  $s_t$  of  $\mathcal{A}$  low.

The sesquilinear product  $\langle \mathcal{X}, \mathcal{Y} \rangle_{\mathcal{A}}$  can be computed without applying  $\mathcal{A}$  to one of the tensors, by interpreting the product as a tensor network and contracting the network. This amounts to a computational complexity of  $O(d(sn^2r + snr^2 + 3s^2r^4 + s^3r^2))$ , where  $\mathcal{X}, \mathcal{Y}$  are in HTD with sizes  $n$  and hierarchical ranks  $r$ , while  $\tilde{\mathcal{A}}$  is in HTD with hierarchical ranks  $s$ .

The composition of two linear operators (i.e., the multiplication of the corresponding matrix representations) in HTD can be calculated in a similar way as the application of a linear operator in HTD.

`y = apply_mat_to_vec(A, x)` returns  $\mathcal{Y} = \mathcal{A}(\mathcal{X})$  for a linear operator  $\mathcal{A}$  in HTD.  
`s = innerprod_mat(x, y, A)` returns  $s = \langle \mathcal{X}, \mathcal{Y} \rangle_{\mathcal{A}}$ .  
`C = apply_mat_to_mat(A, B, p)` returns  $\mathcal{C} = \mathcal{A} \circ \mathcal{B}$  for two linear operators  $\mathcal{A} \in \mathcal{L}(\mathbb{C}^{p_1 \times \dots \times p_d}, \mathbb{C}^{n_1 \times \dots \times n_d})$  and  $\mathcal{B} \in \mathcal{L}(\mathbb{C}^{m_1 \times \dots \times m_d}, \mathbb{C}^{p_1 \times \dots \times p_d})$ .

Note that a truncation of a linear operator  $\mathcal{A}$  to HTD produces a quasi-optimal approximation in the Frobenius norm. Therefore, its effect on the smallest eigenvalues of  $\mathcal{A}$  is likely significant and its use for the direct solution of linear systems questionable. However, this approximation may still be useful in the construction of preconditioners, and there are some notable cases for which an exact representation in HTD of low rank is possible.

*Example 8.1.* A discretized Laplace-like operator of the form

$$A^{(d)} \otimes I \otimes \dots \otimes I + I \otimes A^{(d-1)} \otimes I \otimes \dots \otimes I + \dots + I \otimes \dots \otimes I \otimes A^{(1)} \quad (22)$$

can be represented exactly in HTD with hierarchical rank 2 for any dimension tree  $\mathcal{T}$ :

$$U_t = [\text{vec}(I), \text{vec}(A^{(t)})] \quad \forall t \in \mathcal{L}(\mathcal{T})$$

$$B_t = \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 0 & 1 \\ 0 & 0 \end{bmatrix} \quad \forall t \in \mathcal{N}(\mathcal{T}) \setminus \{t_{\text{root}}\}, \quad B_{t_{\text{root}}} = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \end{bmatrix}.$$

A similar decomposition has been proposed for the TT decomposition in Kazeev and Khoromskij [2012].

## 9. CONCLUSIONS AND ONGOING WORK

The main goal of this work was to provide a convenient way to work with tensors in HTD. Several extensions of this work are possible and currently under consideration.

The most expensive operations of the htucker toolbox are typically calls to level-3 BLAS and LAPACK functions, and hence there is a simple way to gain performance by, for example, linking to multithreaded BLAS. However, in order to address challenging applications that feature high ranks, there is clearly a need for a more advanced implementation fine-tuned to modern high-performance and parallel machines.

One of the most challenging aspects in using low-rank tensor decompositions, such as HTD, is the a priori choice of a suitable decomposition. This also includes the choice of the dimension tree in HTD, which at the moment is ad hoc. The development of rational criteria or effective heuristics for these choices is an important open question.



The work on low-rank tensor decompositions is quickly expanding. As new methods emerge and mature, they will be added to future versions of the htucker toolbox. Candidates include dynamical low-rank methods for HTD [Arnold and Jahnke 2012; Lubich et al. 2013] and low-rank tensor cross-approximation for HTD [Ballani 2012; Ballani et al. 2012]. For both, preliminary implementations are already available; see <http://anchp.epfl.ch/htucker>.

## REFERENCES

- Acar, E., Dunlavy, D. M., and Kolda, T. G. 2011. A scalable optimization approach for fitting canonical tensor decompositions. *J. Chemometr.* 25, 2, 67–86.
- Andersson, C. A. and Bro, R. 2000. The N-Way Toolbox for MATLAB. *Chemometr. Intell. Lab. Syst.* 52, 1, 1–4. <http://www.models.life.ku.dk/nwaytoolbox>.
- Arnold, A. and Jahnke, T. 2012. On the approximation of high-dimensional differential equations in the hierarchical Tucker format. Tech. rep., KIT, Karlsruhe, Germany. [http://www.math.kit.edu/ianm3/~jahnke/seite/jahnke\\_publications/media/arnold-jahnke.pdf](http://www.math.kit.edu/ianm3/~jahnke/seite/jahnke_publications/media/arnold-jahnke.pdf).
- Bader, B. W. and Kolda, T. G. 2006. Algorithm 862: MATLAB tensor classes for fast algorithm prototyping. *ACM Trans. Math. Softw.* 32, 4, 635–653. <http://csmr.ca.sandia.gov/~tgkolda/TensorToolbox/>.
- Bader, B. W. and Kolda, T. G. 2007. Efficient MATLAB computations with sparse and factored tensors. *SIAM J. Sci. Comput.* 30, 1, 205–231.
- Ballani, J. 2012. Fast evaluation of near-field boundary integrals using tensor approximations. Dissertation, Universität Leipzig.
- Ballani, J. and Grasedyck, L. 2013. A projection method to solve linear systems in tensor format. *Numer. Linear Algebra Appl.* 20, 1, 27–43.
- Ballani, J., Grasedyck, L., and Kluge, M. 2012. Black box approximation of tensors in hierarchical tucker format. *Linear Algebra Appl.* 438, 2, 639–657.
- Bauer, B., Carr, L. D., Evertz, H. G., Feiguin, A., Freire, J., et al. 2011. The ALPS project release 2.0: Open source software for strongly correlated systems. *J. Statist. Mech.* 5.
- Carroll, J. and Chang, J.-J. 1970. Analysis of individual differences in multidimensional scaling via an n-way generalization of “Eckart-Young” decomposition. *Psychometrika* 35, 3, 283–319.
- De Lathauwer, L., De Moor, B., and Vandewalle, J. 2000. A multilinear singular value decomposition. *SIAM J. Matrix Anal. Appl.* 21, 4, 1253–1278.
- Espig, M. and Hackbusch, W. 2012. A regularized Newton method for the efficient approximation of tensors represented in the canonical tensor format. *Numer. Math.* 122, 3, 489–525.
- Espig, M., Hackbusch, W., Handschuh, S., and Schneider, R. 2011. Optimization problems in contracted tensor networks. *Comput. Visual. Sci.* 14, 271–285.
- Espig, M., Hackbusch, W., Rohwedder, T., and Schneider, R. 2012a. Variational calculus with sums of elementary tensors of fixed rank. *Numer. Math.* 122, 3, 469–488.
- Espig, M., Schuster, M., Killaitis, A., Waldren, N., Wähnert, P., Handschuh, S., and Auer, H. 2012b. Tensor-Calculus library. <http://gitorious.org/tensorcalculus>.
- Golub, G. H. and Van Loan, C. F. 1996. *Matrix Computations* 3rd Ed. Johns Hopkins University Press, Baltimore, MD.
- Grasedyck, L. 2004. Existence and computation of low Kronecker-rank approximations for large linear systems of tensor product structure. *Comput.* 72, 3–4, 247–265.
- Grasedyck, L. 2010a. Hierarchical low rank approximation of tensors and multivariate functions. In *Lecture Notes of Zurich Summer School on Sparse Tensor Discretizations of High-Dimensional Problems*.
- Grasedyck, L. 2010b. Hierarchical singular value decomposition of tensors. *SIAM J. Matrix Anal. Appl.* 31, 4, 2029–2054.
- Hackbusch, W. 2012. *Tensor Spaces and Numerical Tensor Calculus*. Series in Computational Mathematics, vol. 42, Springer.
- Hackbusch, W. and Kuhn, S. 2009. A new scheme for the tensor representation. *J. Fourier Anal. Appl.* 15, 5, 706–722.
- Hackbusch, W., Khoromskij, B. N., Sauter, S. A., and Tyrtshnikov, E. E. 2012. Use of tensor formats in elliptic eigenvalue problems. *Numer. Linear Algebra Appl.* 19, 1, 133–151.
- Harshman, R. A. 1970. Foundations of the PARAFAC procedure: Models and conditions for an “explanatory” multi-modal factor analysis. *UCLA Working Papers Phonet.* 16, 1–84.
- Holtz, S., Rohwedder, T., and Schneider, R. 2012. The alternating linear scheme for tensor optimization in the tensor train format. *SIAM J. Sci. Comput.* 34, 2, A683–A713.

- Huckle, T., Waldherr, K., and Schulte-Herbrüggen, T. 2013. Computations in quantum tensor networks. *Linear Algebra Appl.* 438, 2, 750–781.
- Kazeev, V. A. and Khoromskij, B. N. 2012. Low-rank explicit QTT representation of Laplace operator and its inverse. *SIAM J. Matrix Anal. Appl.* 33, 3, 742–758.
- Khoromskij, B. N. and Oseledets, I. V. 2010. Quantics-TT collocation approximation of parameter-dependent and stochastic elliptic PDES. *Comput. Methods Appl. Math.* 10, 4, 376–394.
- Khoromskij, B. N. and Oseledets, I. V. 2011. QTT approximation of elliptic solution operators in higher dimensions. *Russian J. Numer. Anal. Math. Model.* 26, 3, 303–322.
- Khoromskij, B. N. and Schwab, C. 2011. Tensor-structured Galerkin approximation of parametric and stochastic elliptic PDES. *SIAM J. Sci. Comput.* 33, 1, 364–385.
- Koch, O. and Lubich, C. 2010. Dynamical tensor approximation. *SIAM J. Matrix Anal. Appl.* 31, 5, 2360–2375.
- Kolda, T. G. and Bader, B. W. 2009. Tensor decompositions and applications. *SIAM Rev.* 51, 3, 455–500.
- Kressner, D. and Tobler, C. 2010. Krylov subspace methods for linear systems with tensor product structure. *SIAM J. Matrix Anal. Appl.* 31, 4, 1688–1714.
- Kressner, D. and Tobler, C. 2011. Low-rank tensor Krylov subspace methods for parametrized linear systems. *SIAM J. Matrix Anal. Appl.* 32, 4, 1288–1316.
- Kressner, D. and Tobler, C. 2012. htucker – A MATLAB Toolbox for tensors in hierarchical Tucker format. Tech. rep. 2012-02, MATHICSE, EPF Lausanne. (Extended version at <http://anchp.epfl.ch/htucker>.)
- Lubich, C., Rohwedder, T., Schneider, R., and Vandereycken, B. 2013. Dynamical approximation of hierarchical Tucker and tensor-train tensors. *SIAM J. Matrix Anal. Appl.* 34, 2, 470–494.
- Oseledets, I. V. 2011a. MATLAB TT-Toolbox version 2.1. <http://spring.inm.ras.ru/osel/?pageid=24>.
- Oseledets, I. V. 2011b. Tensor-train decomposition. *SIAM J. Sci. Comput.* 33, 5, 2295–2317.
- Oseledets, I. V. and Khoromskij, B. N. 2010. DMRG+QTT approach to high-dimensional quantum molecular dynamics. Preprint 69/2010. Max-Planck-Institut für Mathematik in den Naturwissenschaften.
- Schollwöck, U. 2011. The density-matrix renormalization group in the age of matrix product states. *Ann. Phys.* 326, 1, 96–192.
- Tobler, C. 2012. Low-rank tensor methods for linear systems and eigenvalue problems. Ph.D. dissertation, ETH Zurich, Switzerland.
- Tucker, L. 1966. Some mathematical notes on three-mode factor analysis. *Psychometrika* 31, 3, 279–311.

Received February 2012; revised September 2013; accepted October 2013